



# PyMC's Big Adventure

*Lessons Learned from the Development of Open-source  
Software for Probabilistic Programming*

Chris Fonnesbeck, Vanderbilt University Medical Center

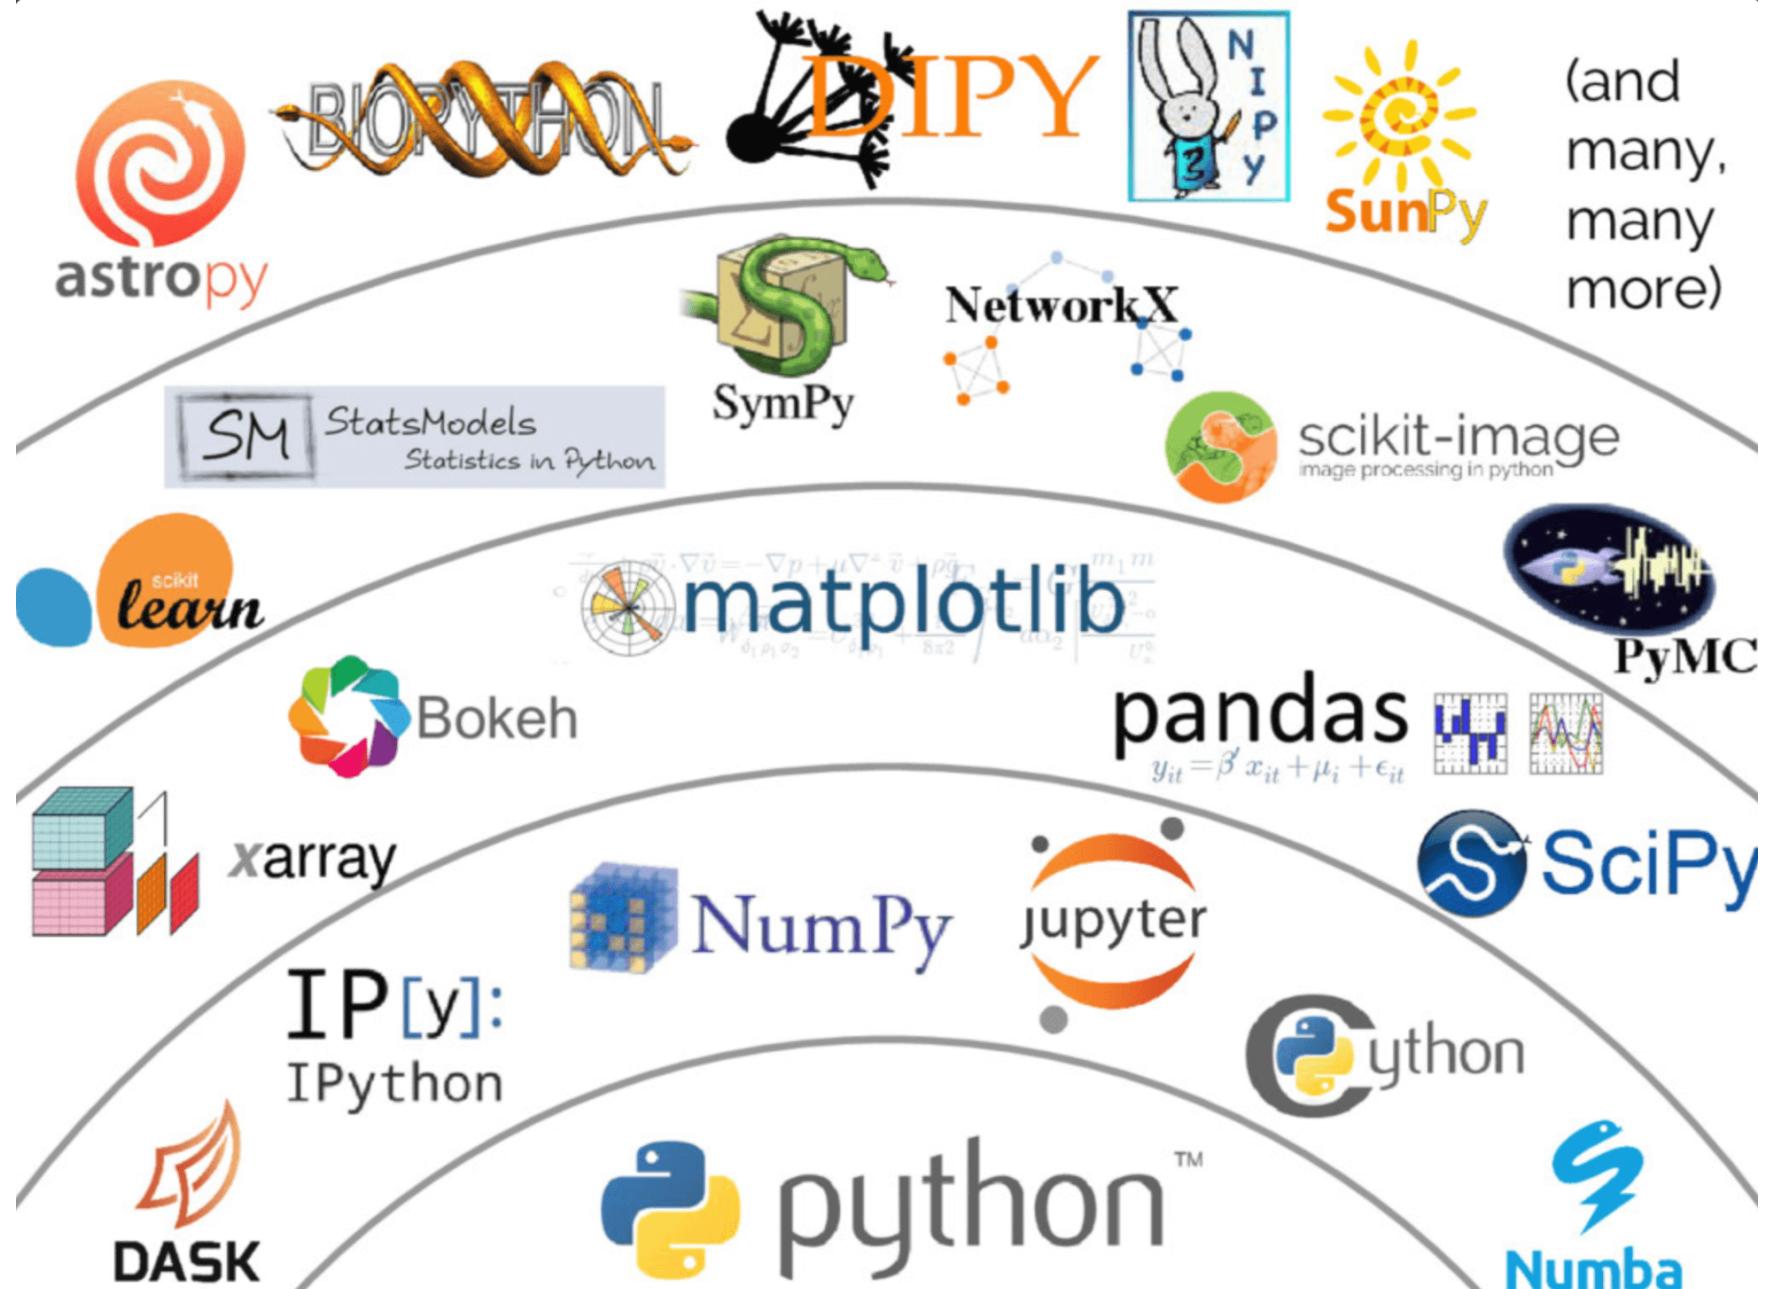
Machine Learning Open Source Software 2018: Sustainable communities



- started in 2003
- **probabilistic programming** framework
- based on Theano
- implements *next generation* Bayesian inference methods
- >100 contributors

≡

Salvatier, Wiecki and Fonnesbeck (2016)



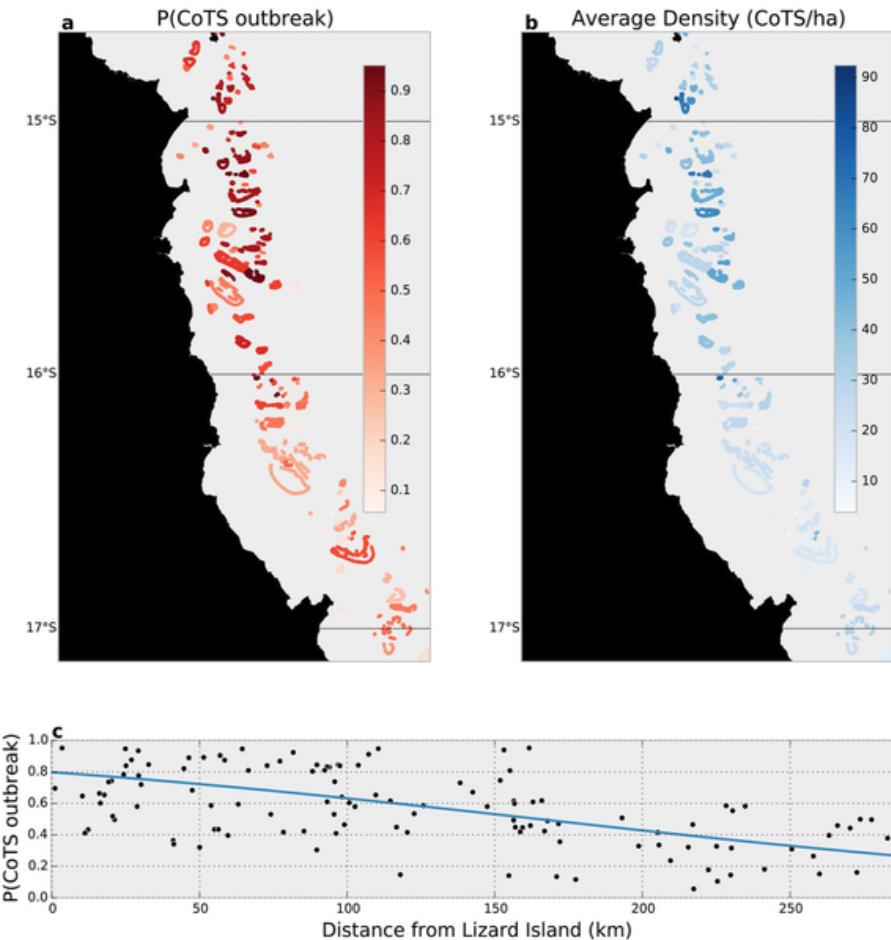
≡



# PyMC3 in the Wild

- Astronomy
- Biostatistics
- Ecology
- Chemistry
- Finance
- Economics
- Sabermetrics
- Marketing
- Actuarial Science

Joint estimation of crown of thorns *Acanthaster planci* densities on the Great Barrier Reef



MacNeil et al. 2016





```
In [19]: import pymc3 as pm

with pm.Model() as framing_model:

    # Baseline probability
    μ = pm.Normal('μ', 2, sd=5)
    σ = pm.Exponential('σ', 1)
    z = pm.Normal('z', 0, 1, shape=n_catchers)
    # Non-centered parameterization
    θ = pm.Deterministic('θ', μ + z*σ)

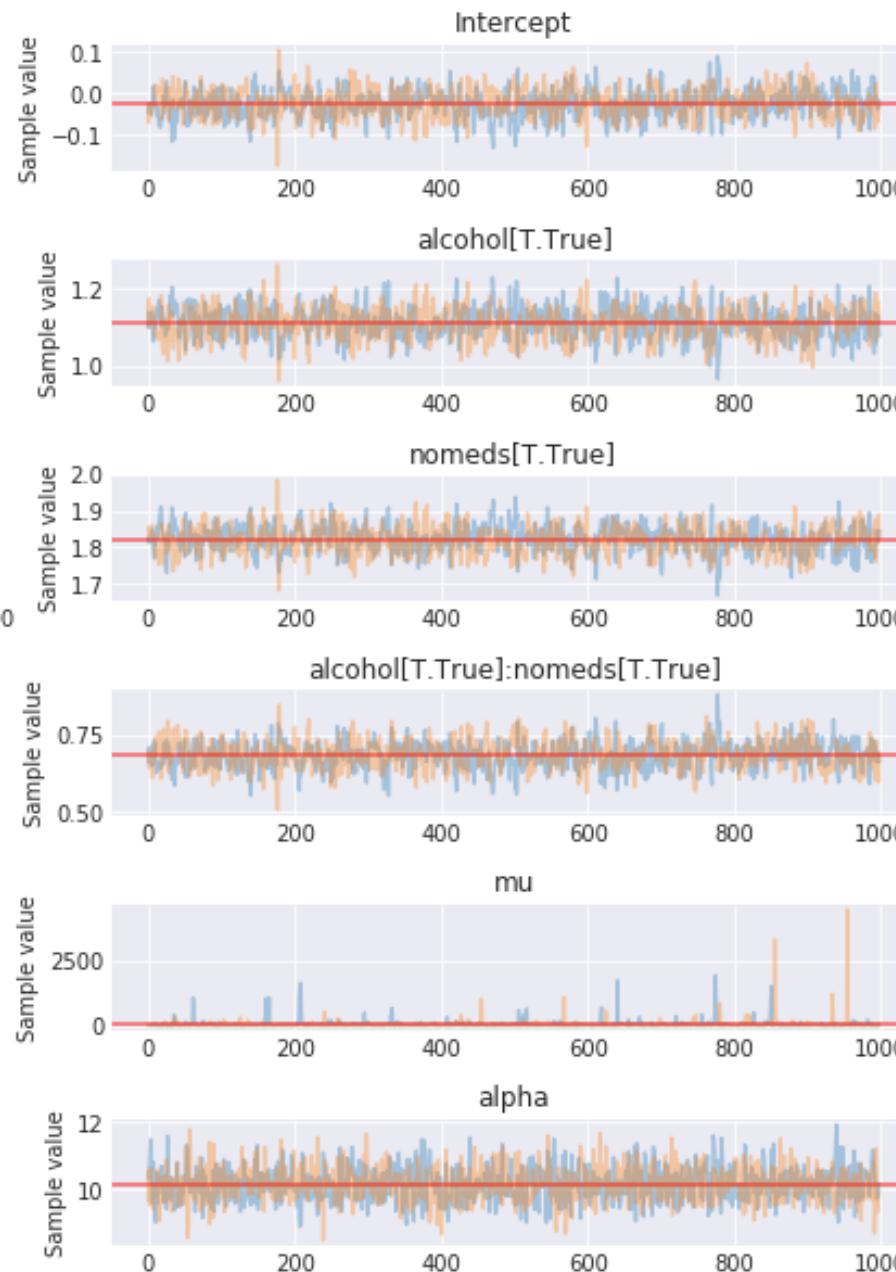
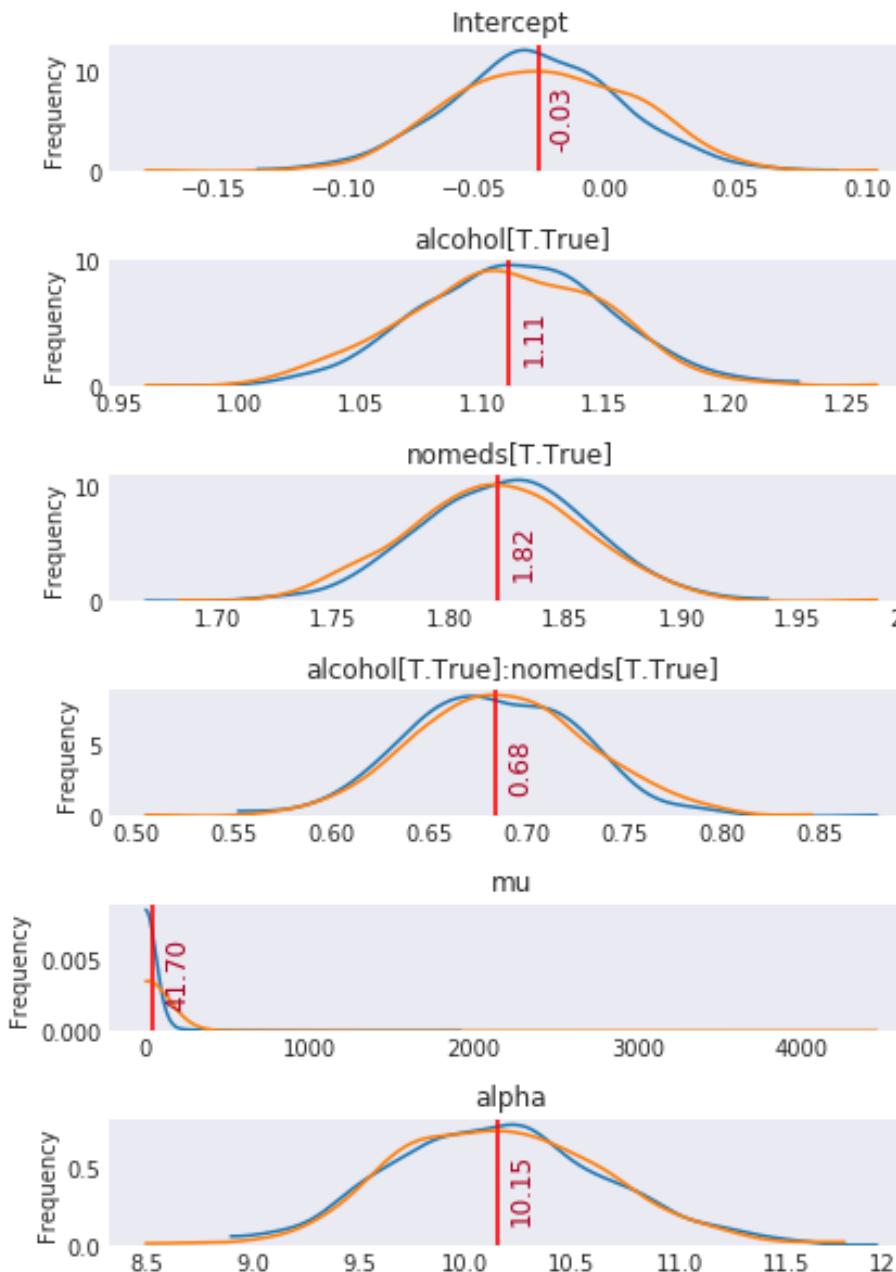
    # Called strike rate as function of distance to strike zone
    p = pm.math.invlogit(θ[catcher_ind])

    called_strike = pm.Bernoulli('called_strike', p, observed=strike)
```

```
In [20]: with framing_model:
    samples = pm.sample(1000, tune=1000, njobs=2)
```

```
INFO:pymc3:Auto-assigning NUTS sampler...
INFO:pymc3:Initializing NUTS using jitter+adapt_diag...
INFO:pymc3:Multiprocess sampling (2 chains in 2 jobs)
INFO:pymc3:NUTS: [z, σ, μ]
Sampling 2 chains: 100%|██████████| 4000/4000 [01:49<00:00, 58.81draws/s]
```







# Why PyMC?

Bayesian modeling for applied users ☺



```
Error: object "Royce" not found
> source("C:/garyking/king_england_royce.R")
Read 36 items
Read 36 items
> model
$par
[1] 1.14217375 -0.0458708 -1.603315

$value
[1] 6118.279

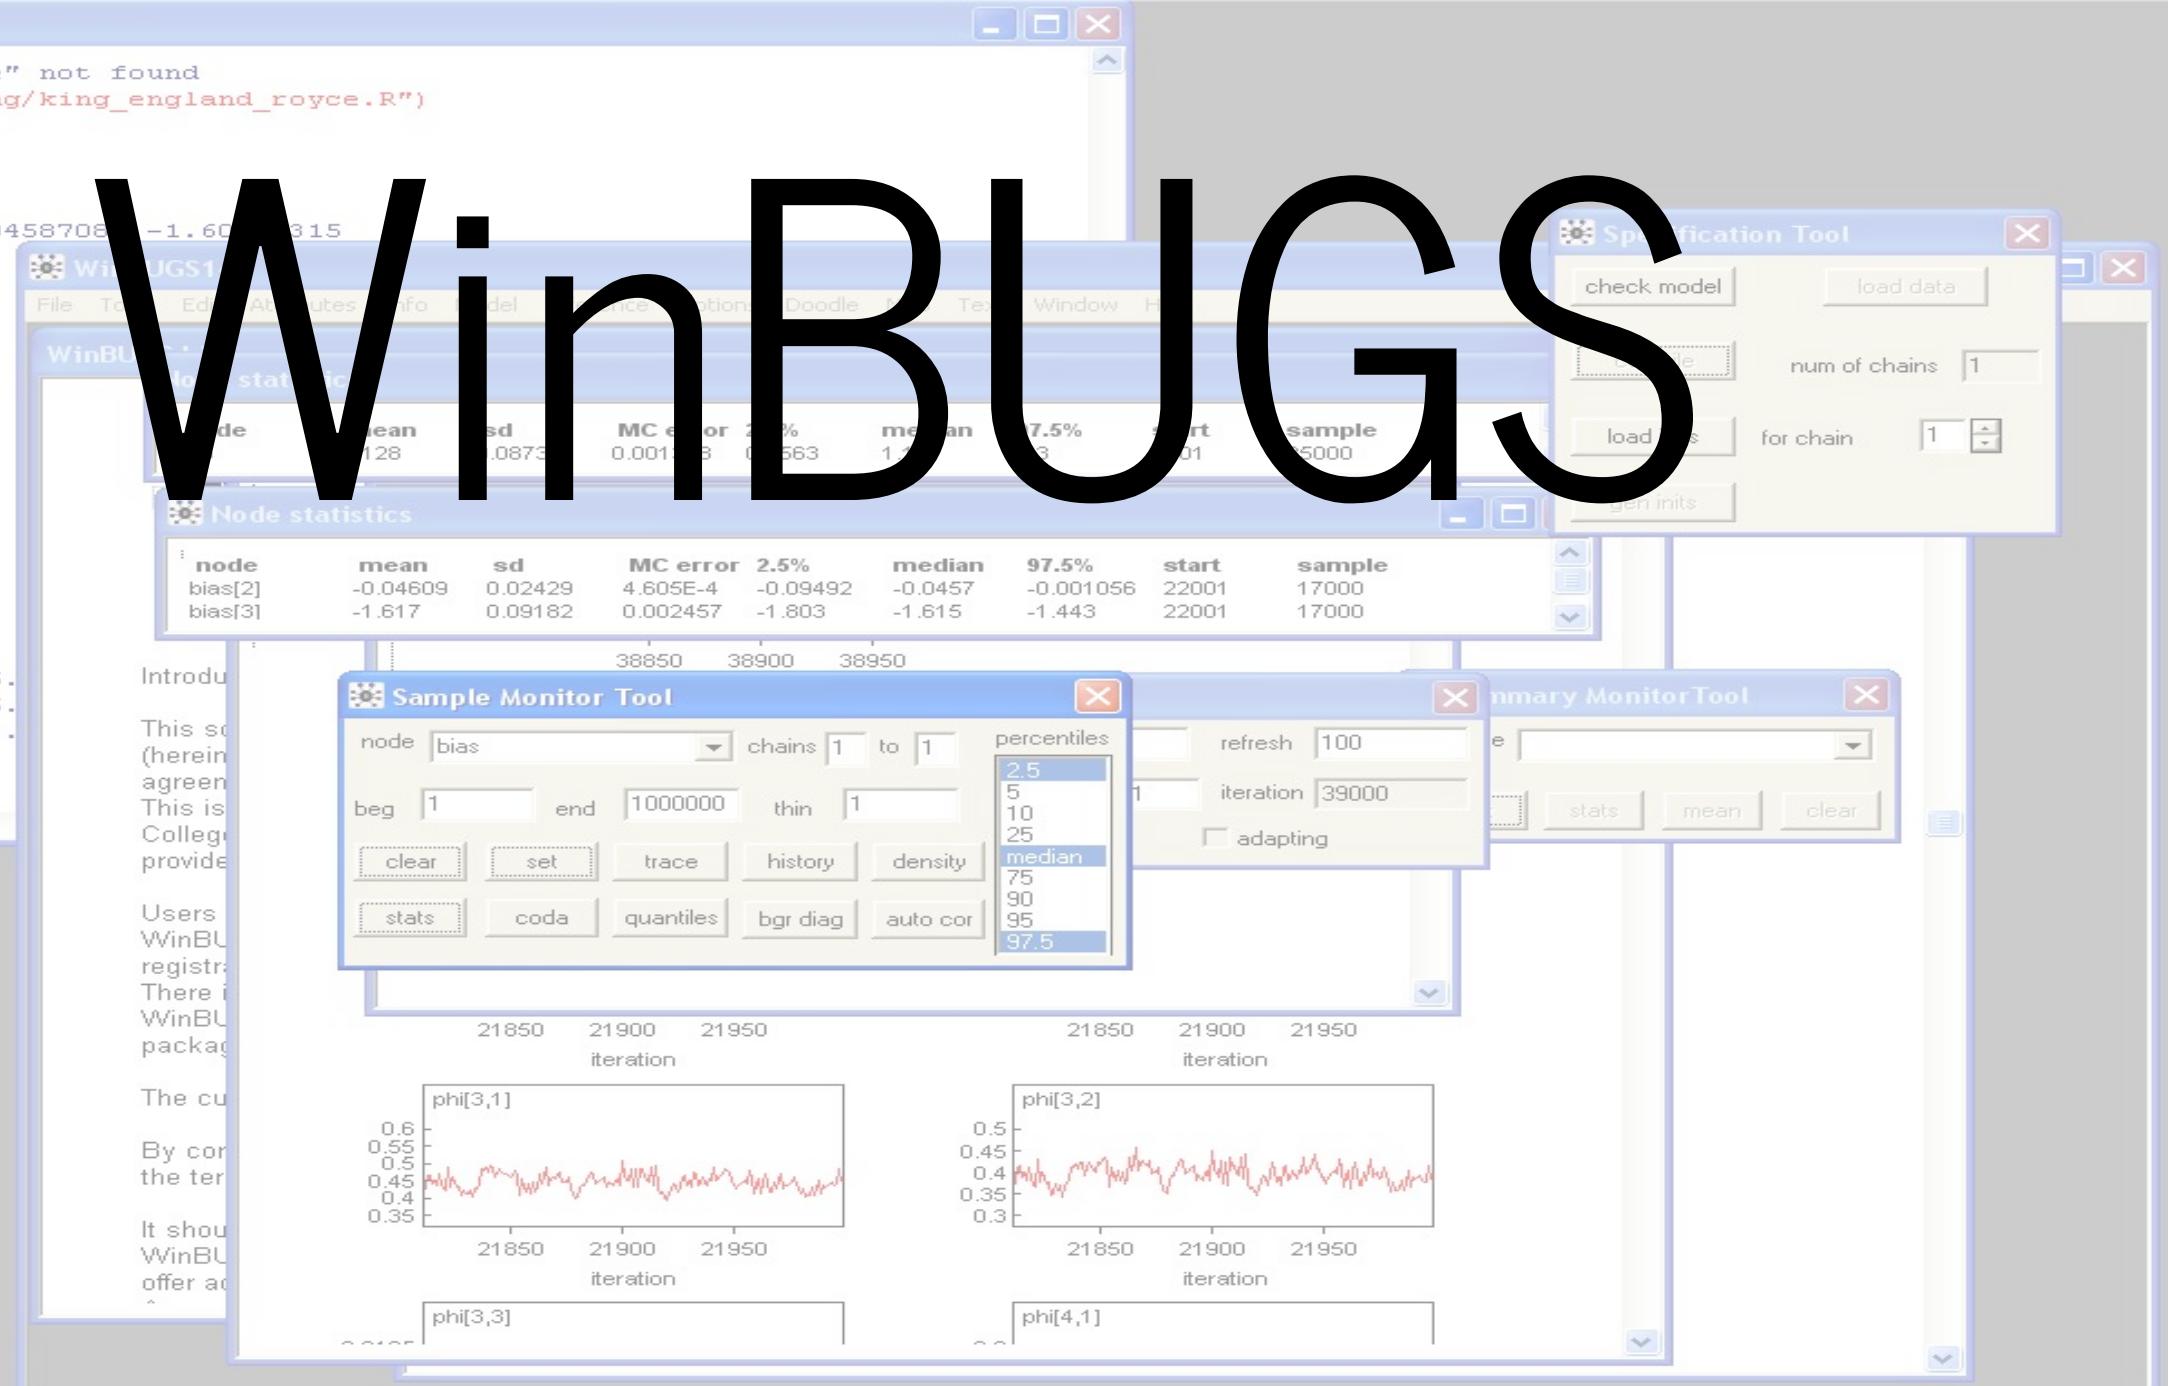
$counts
function gradient
 264      NA

$convergence
[1] 0

$message
NULL

$hessian
[,1]
[1,] 254.92549 -45.
[2,] -45.61907 1866.
[3,] -171.28432 -112.

> 
```



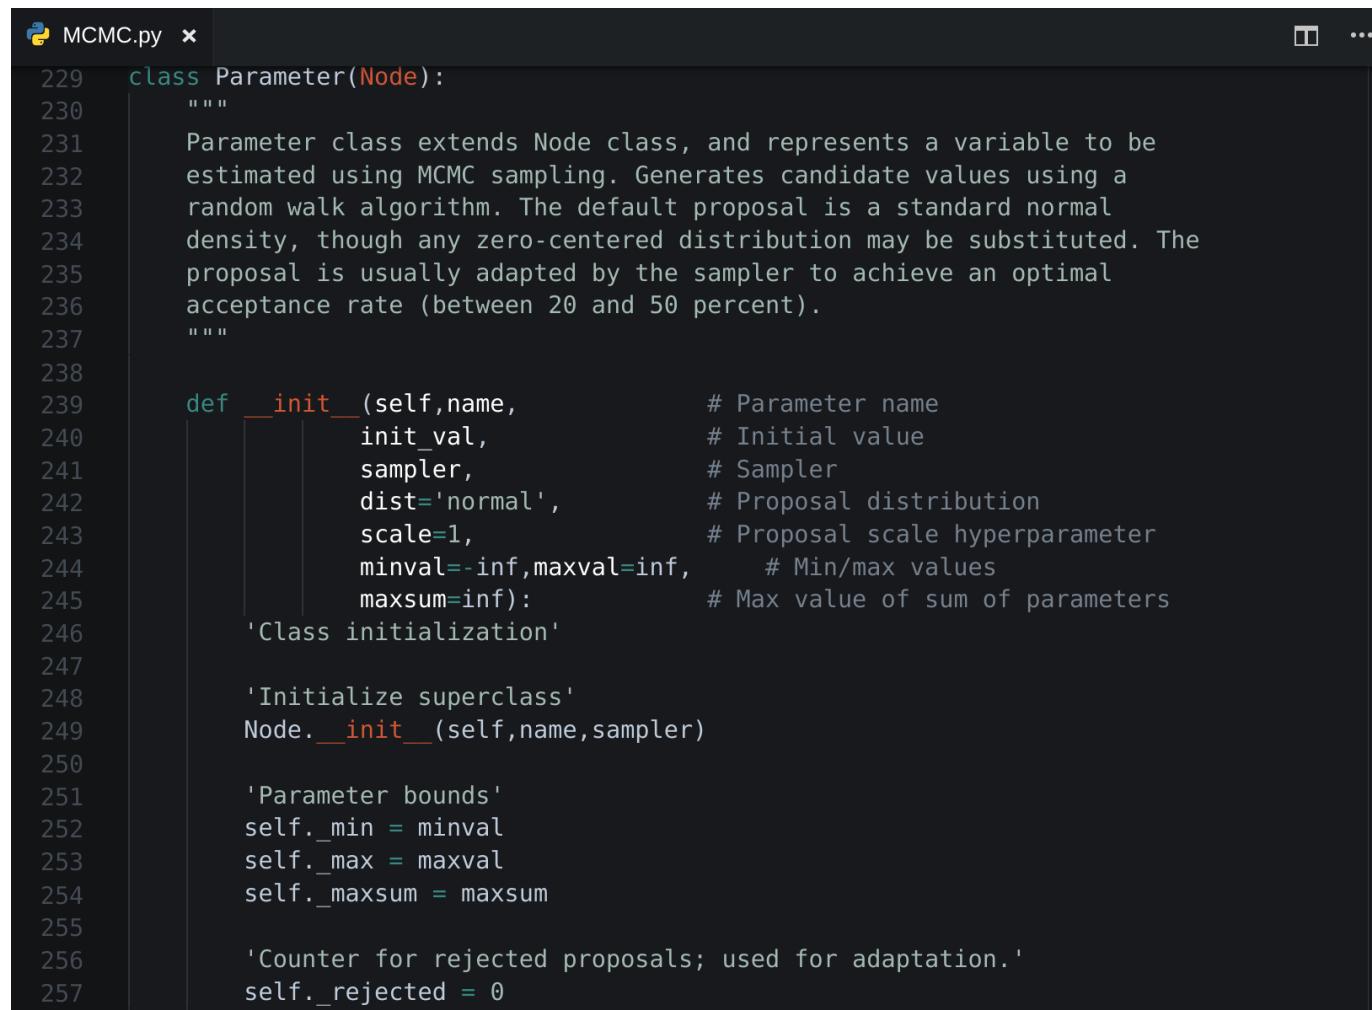


```
model {  
    for (j in 1:J){  
        y[j] ~ dnorm (theta[j], tau.y[j])  
        theta[j] ~ dnorm (mu.theta, tau.theta)  
        tau.y[j] <- pow(sigma.y[j], -2)  
    }  
    mu.theta ~ dnorm (0.0, 1.0E-6)  
    tau.theta <- pow(sigma.theta, -2)  
    sigma.theta ~ dunif (0, 1000)  
}
```





# PyMC 1.0



The screenshot shows a code editor window with the file 'MCMC.py' open. The code defines a class 'Parameter' that extends the 'Node' class. The class has a detailed docstring explaining its purpose: generating candidate values using a random walk algorithm. It includes parameters for proposal distribution ('dist'), proposal scale hyperparameter ('scale'), and min/max values ('minval', 'maxval', 'maxsum'). The constructor initializes the parameter name, sampler, and bounds. It also initializes a counter for rejected proposals ('\_rejected').

```
229  class Parameter(Node):
230      """
231          Parameter class extends Node class, and represents a variable to be
232          estimated using MCMC sampling. Generates candidate values using a
233          random walk algorithm. The default proposal is a standard normal
234          density, though any zero-centered distribution may be substituted. The
235          proposal is usually adapted by the sampler to achieve an optimal
236          acceptance rate (between 20 and 50 percent).
237      """
238
239      def __init__(self, name,
240                  init_val, # Initial value
241                  sampler, # Sampler
242                  dist='normal', # Proposal distribution
243                  scale=1, # Proposal scale hyperparameter
244                  minval=-inf, maxval=inf, # Min/max values
245                  maxsum=inf): # Max value of sum of parameters
246          'Class initialization'
247
248          'Initialize superclass'
249          Node.__init__(self, name, sampler)
250
251          'Parameter bounds'
252          self._min = minval
253          self._max = maxval
254          self._maxsum = maxsum
255
256          'Counter for rejected proposals; used for adaptation.'
257          self._rejected = 0
```





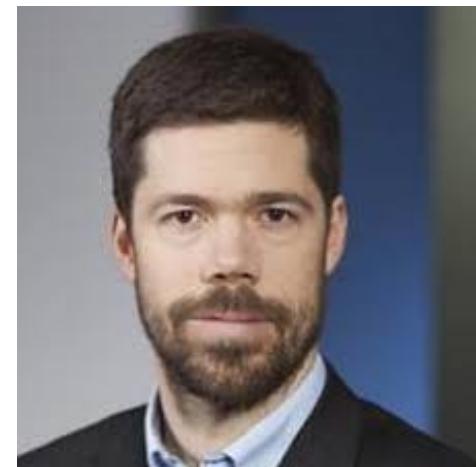
# And Then There Were Three ...



Anand Patil



Me



David Huard

≡



# PyMC2

- Suite of well-documented statistical distributions.
- NumPy-based
- Gaussian processes module
- Robust sampling loops
- Suite of convergence diagnostics
- Extensible: custom step methods and unusual probability distributions
- f2py FORTRAN extensions

≡

Patil, A., D. Huard, and C.J. Fonnesbeck. 2010. Journal of Statistical Software 35 [4].



## pymc-devs / pymc

Code

Issues 13

Pull requests 1

PyMC: Bayesian Stochastic Modelling in Python (for  
[pymc-devs.github.com/pymc/](https://pymc-devs.github.com/pymc/))

Manage topics

● Fortran 75.9%

● Python 16.2%

Branch: master ▾

New pull request





# Hamiltonian Monte Carlo

Uses a *physical analogy* of a frictionless particle moving on a hyper-surface

Requires an *auxiliary variable* to be specified

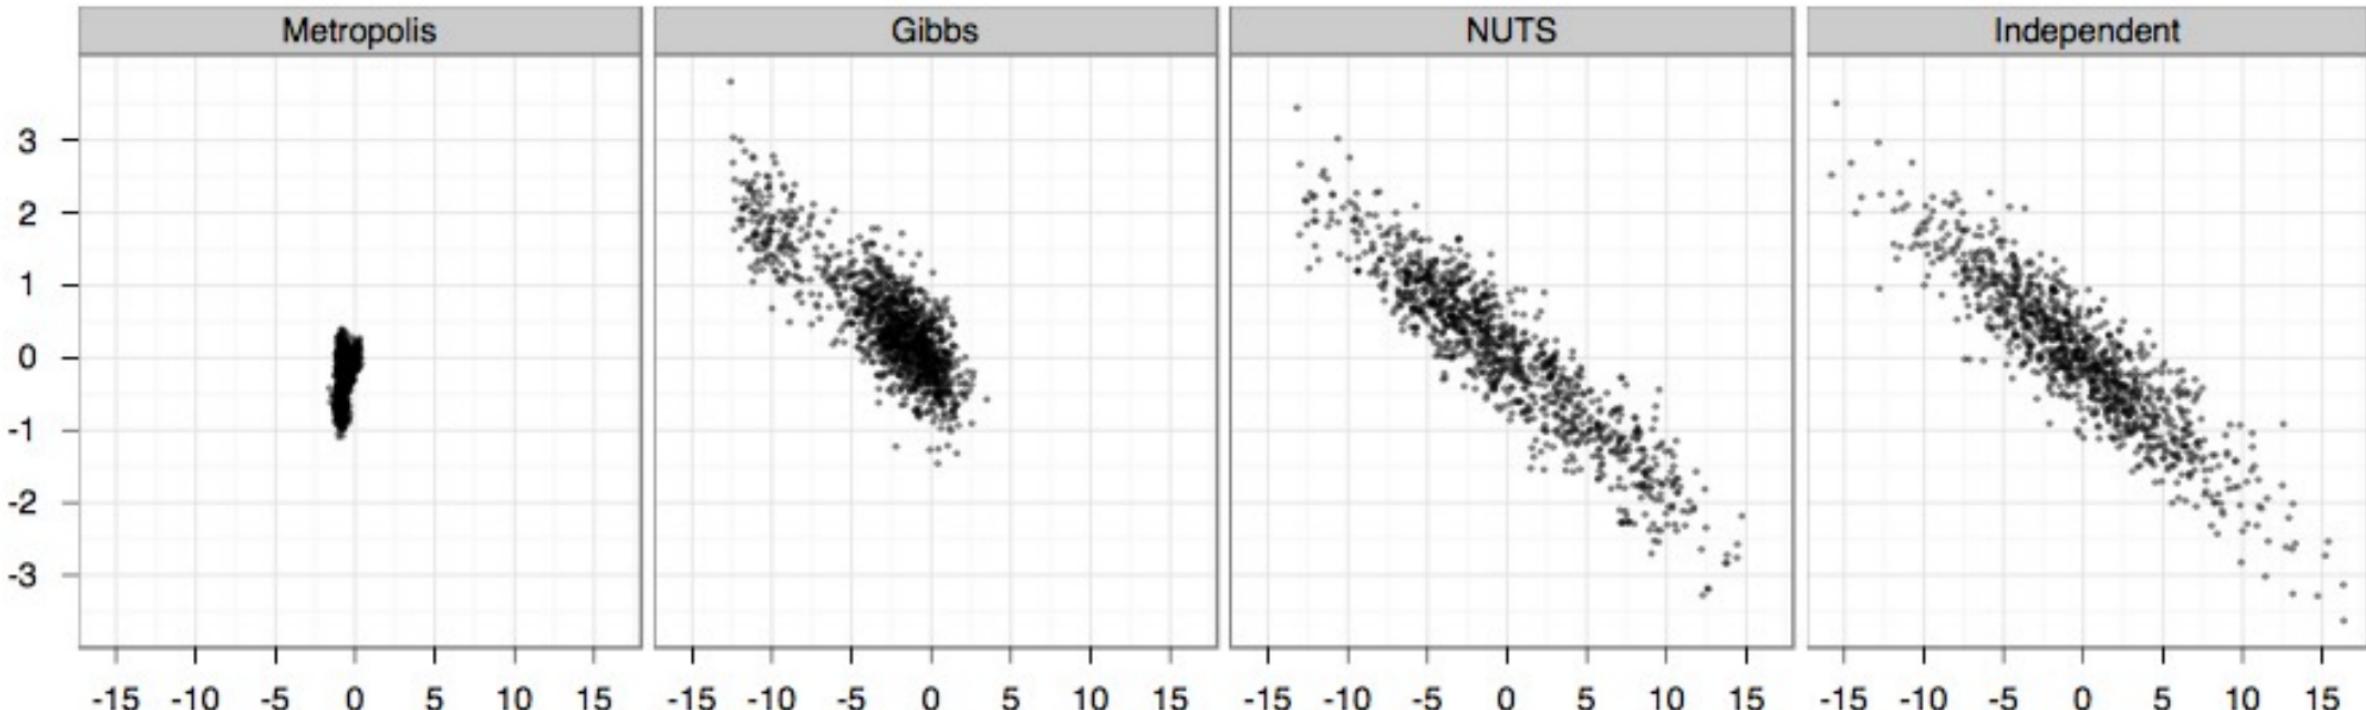
- position (unknown variable value)
- momentum (auxiliary)

$$\mathcal{H}(s, \phi) = E(s) + K(\phi) = E(s) + \frac{1}{2} \left( \sum_i \right) \phi_i^2$$

≡



# Efficient MCMC sampling



Hoffman and Gelman, 2014



"I'm not sure why this approach seems neglected. It might be that research incentives don't reward such generally applicable research, or that MCMC researchers do not see how simplified MCMC could dramatically improve the productivity of statistics, or perhaps researchers haven't realized how automatic differentiation can democratize these algorithms."



John Salvatier



## Calculating Gradients in Theano

```
>>> from theano import function, tensor as tt
>>> x = tt.dmatrix('x')
>>> s = tt.sum(1 / (1 + tt.exp(-x)))
>>> gs = tt.grad(s, x)
>>> dlogistic = function([x], gs)
>>> dlogistic([[3, -1], [0, 2]])
array([[ 0.04517666,  0.19661193],
       [ 0.25        ,  0.10499359]])
```





# PyMC Devs, The Next Generation



John Salvatier



Me



Thomas Wiecki

≡



# Regression Model

$$\beta_{0,i} \sim N(0, 10000)$$

$$\beta_1 \sim N(0, 10000)$$

$$\sigma \sim C^+(5)$$

$$\theta_i = \beta_{0,i} + \beta_1 x_i$$

$$y_i \sim N(\theta_i, \sigma)$$

≡



# Regression Model

```
with Model() as unpooled_model:  
  
    β₀ = Normal('β₀', 0, sd=1e5, shape=counties)  
    β₁ = Normal('β₁', 0, sd=1e5)  
    σ = HalfCauchy('σ', 5)  
  
    θ = β₀[county] + β₁*floor  
  
    y = Normal('y', θ, sd=σ, observed=log_radon)
```

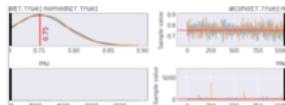




Bayes' Theorem for 2 variables:  
let  $E, F$  be events  
with  $P(E) \neq 0$  and  $P(F) \neq 0$

Then

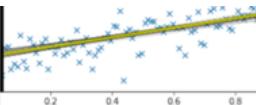
$$P(F|E) = \frac{P(E|F)P(F)}{P(E|F)P(F) + P(E|\bar{F})P(\bar{F})}$$
$$P(E|F) = \frac{P(EF)}{P(F)} \rightarrow P(E|F)P(F) = P(EF)$$
$$P(F|E) = \frac{P(EF)}{P(E)} \rightarrow P(F|E)P(E) = P(EF)$$
$$P(E|F)P(F) = P(F|E)P(E)$$
$$P(E|F) = \frac{P(F|E)P(E)}{P(F)}$$

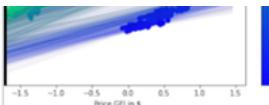
GLM: Poisson  
Regression



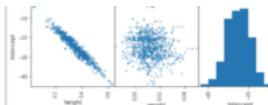
GLM: Robust  
Regression with  
Outlier  
Detection



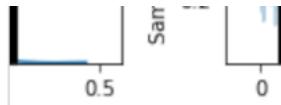
GLM: Robust  
Linear  
Regression



Rolling  
Regression

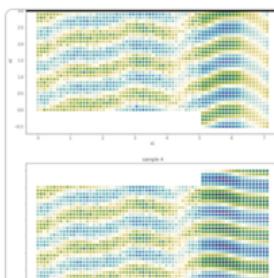


GLM: Linear  
Regression

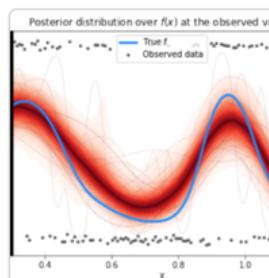


Hierarchical  
Partial Pooling

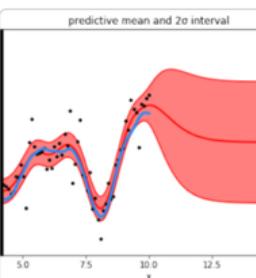
## Gaussian Processes



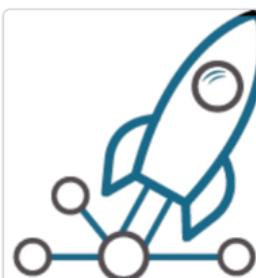
Kronecker  
Structured  
Covariances



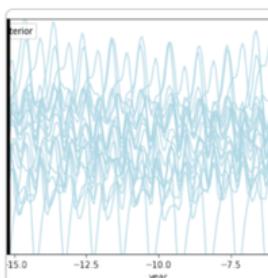
Latent Variable  
Implementation



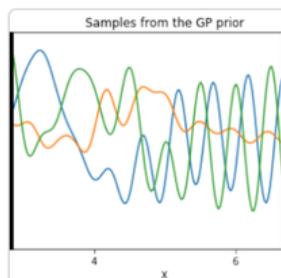
Marginal  
Likelihood  
Implementation



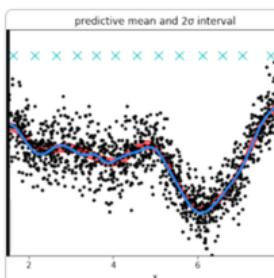
Example: CO<sub>2</sub> at  
Mauna Loa



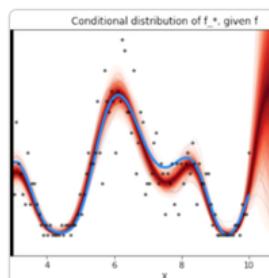
Example: Mauna  
Loa CO<sub>2</sub>\_2\$  
continued



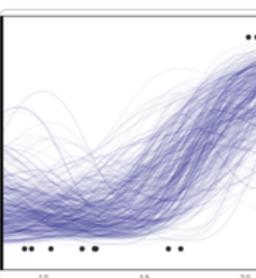
Mean and  
Covariance  
Functions



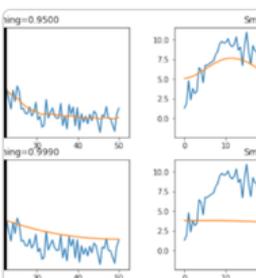
Sparse  
Approximations



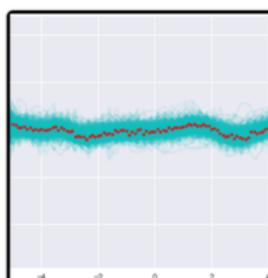
Student-t  
Process



Gaussian Process  
Regression and  
Classification



Gaussian Process  
(GP) smoothing



Gaussian  
Processes



# MCMC Sampling

```
with pitch_framing_model:  
    trace = pm.sample(1000, tune=1000)  
  
Auto-assigning NUTS sampler...  
Initializing NUTS using jitter+adapt_diag...  
Multiprocess sampling (2 chains in 2 jobs)  
NUTS: [σ_log__, ψ, μ]  
100%|██████████| 2000/2000 [00:42<00:00, 47.42it/s]
```





# Variational Inference

```
with neural_network:  
    approx = pm.fit(n=30000, method='fullrank_advi')
```

```
Average Loss = 1,207.6: 100%|██████████| 30000/30000 [27:41  
Finished [100%]: Average Loss = 1,206.2
```





## ★ MILA and the future of Theano



Pascal Lamblin

Dear users and developers,

After almost ten years of development, we have the regret to announce that we will put an end to our Theano development after the 1.0 release, which is due in the next few weeks. We will continue minimal maintenance to keep it working for one year, but we will stop actively implementing new features. Theano will continue to be available afterwards, as per our engagement towards open source software, but MILA does not commit to spend time on maintenance or support after that time frame.

The software ecosystem supporting deep learning research has been evolving quickly, and has now reached a healthy state: open-source software is the norm; a variety of frameworks are available, satisfying needs spanning from exploring novel ideas to deploying them into production; and strong industrial players are backing different software stacks in a stimulating competition.

We are proud that most of the innovations Theano introduced across the years have now been adopted and perfected by other frameworks. Being able to express models as mathematical expressions, rewriting







**Chris Fonnesbeck**

@fonnesbeck



The #PyMC3 devs are frequently asked about what's going to happen to the project now that Theano is winding down. Here is a summary of where things are headed.



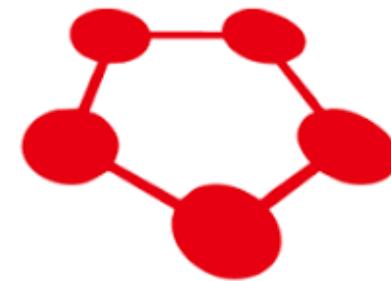
### **Theano, TensorFlow and the Future of PyMC – PyMC Devel...**

PyMC3 is an open-source library for Bayesian statistical modeling and inference in Python, implementing gradient-based Markov chain Monte...

[medium.com](https://medium.com/@fonnesbeck/theano-tensorflow-and-the-future-of-pymc-pymc-devel-175a2a2a2a2)



2:50 PM - 17 May 2018



Chainer

PYTORCH







# NUMFOCUS

OPEN CODE = BETTER SCIENCE

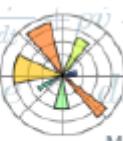
## SPONSORED PROJECTS



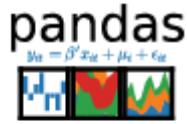
NumPy



IP[y]:  
IPython



Matplotlib



pandas



interact



Stan



將軍



open  
Journals



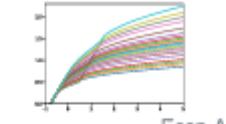
SymPy



julia



PyTables



FENICS  
PROJECT



Econ-ARK



AstroPy



yt



sunpy

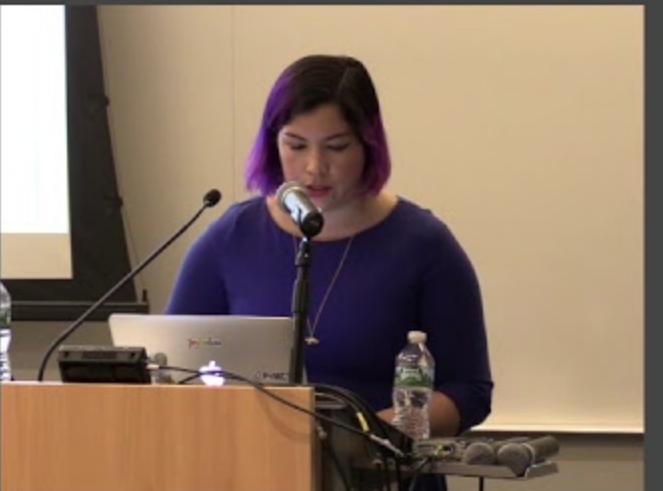




Google

SUMMER  
of CODE





```
In [19]: summary(trace)

alpha:

```

Mean	SD	MC Error	95% HPD interval
3.002	0.220	0.003	[2.582, 3.440]

```
Posterior quantiles:
  2.5          25          50          75          97.5
 |-----|=====|=====|=====|-----|
  2.579       2.851       3.000       3.150       3.440
```

```
beta:

```

Mean	SD	MC Error	95% HPD interval
3.968	0.238	0.003	[3.501, 4.423]

```
Posterior quantiles:
  2.5          25          50          75          97.5
 |-----|=====|=====|=====|-----|
  3.502       3.809       3.966       4.129       4.424
```

```
s:

```

Mean	SD	MC Error	95% HPD interval
2.021	0.149	0.002	[1.734, 2.306]

```
Posterior quantiles:
  2.5          25          50          75          97.5
 |-----|=====|=====|=====|-----|
  1.743       1.918       2.014       2.117       2.322
```





# Interaction and Communication

- Monthly lab meetings
- Journal club
- Face-to-face meetings
- Slack channels
- Discourse site

≡

